

Kafka Consumers



Axel Sirota

AI and Cloud Consultant

@AxelSirota


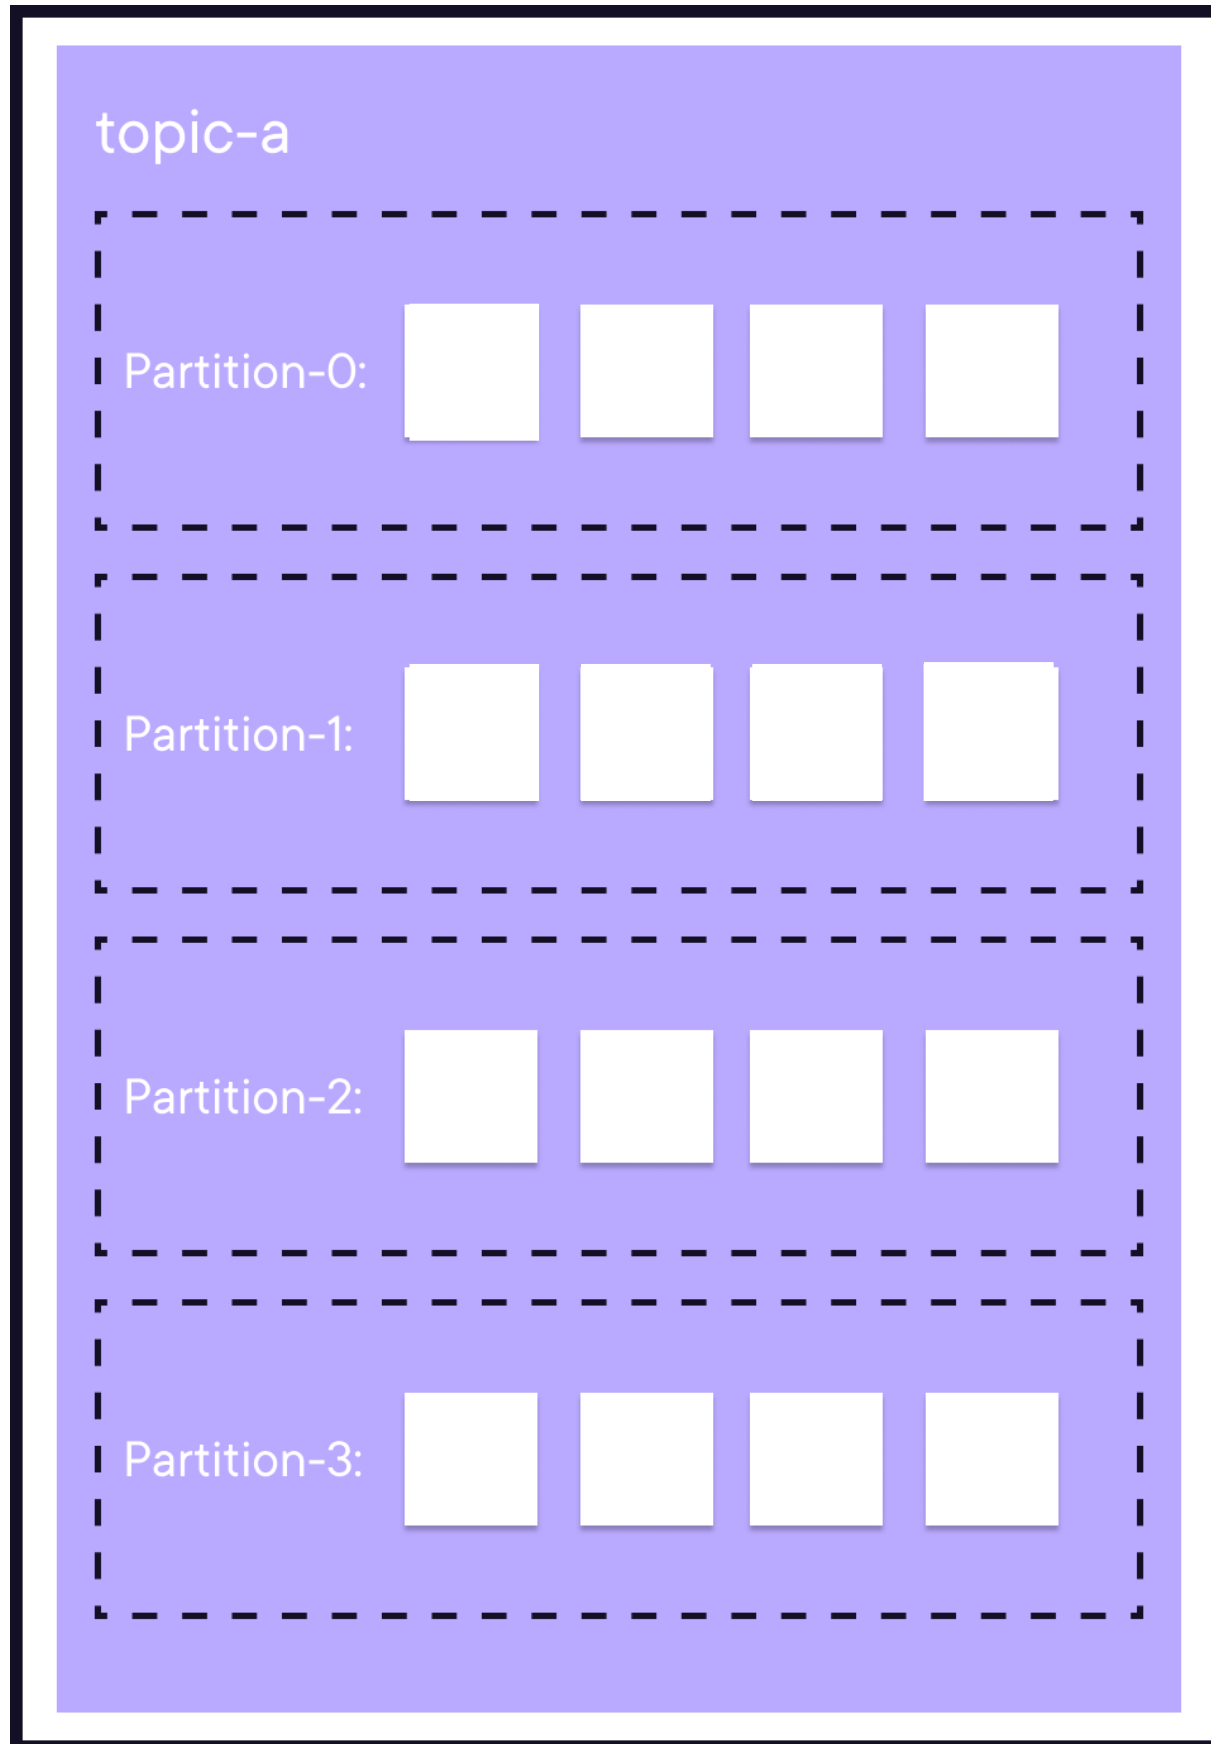


Kafka Consumer Group




**Consumers are
typically done as a
group**





Kafka Consumer Group



**Consumers are
typically done as a
group**


**A single consumer will
end up inefficient with
large amounts of data**




Kafka Consumer Group



Consumers are typically done as a group



A single consumer will end up inefficient with large amounts of data





A consumer may never catch up

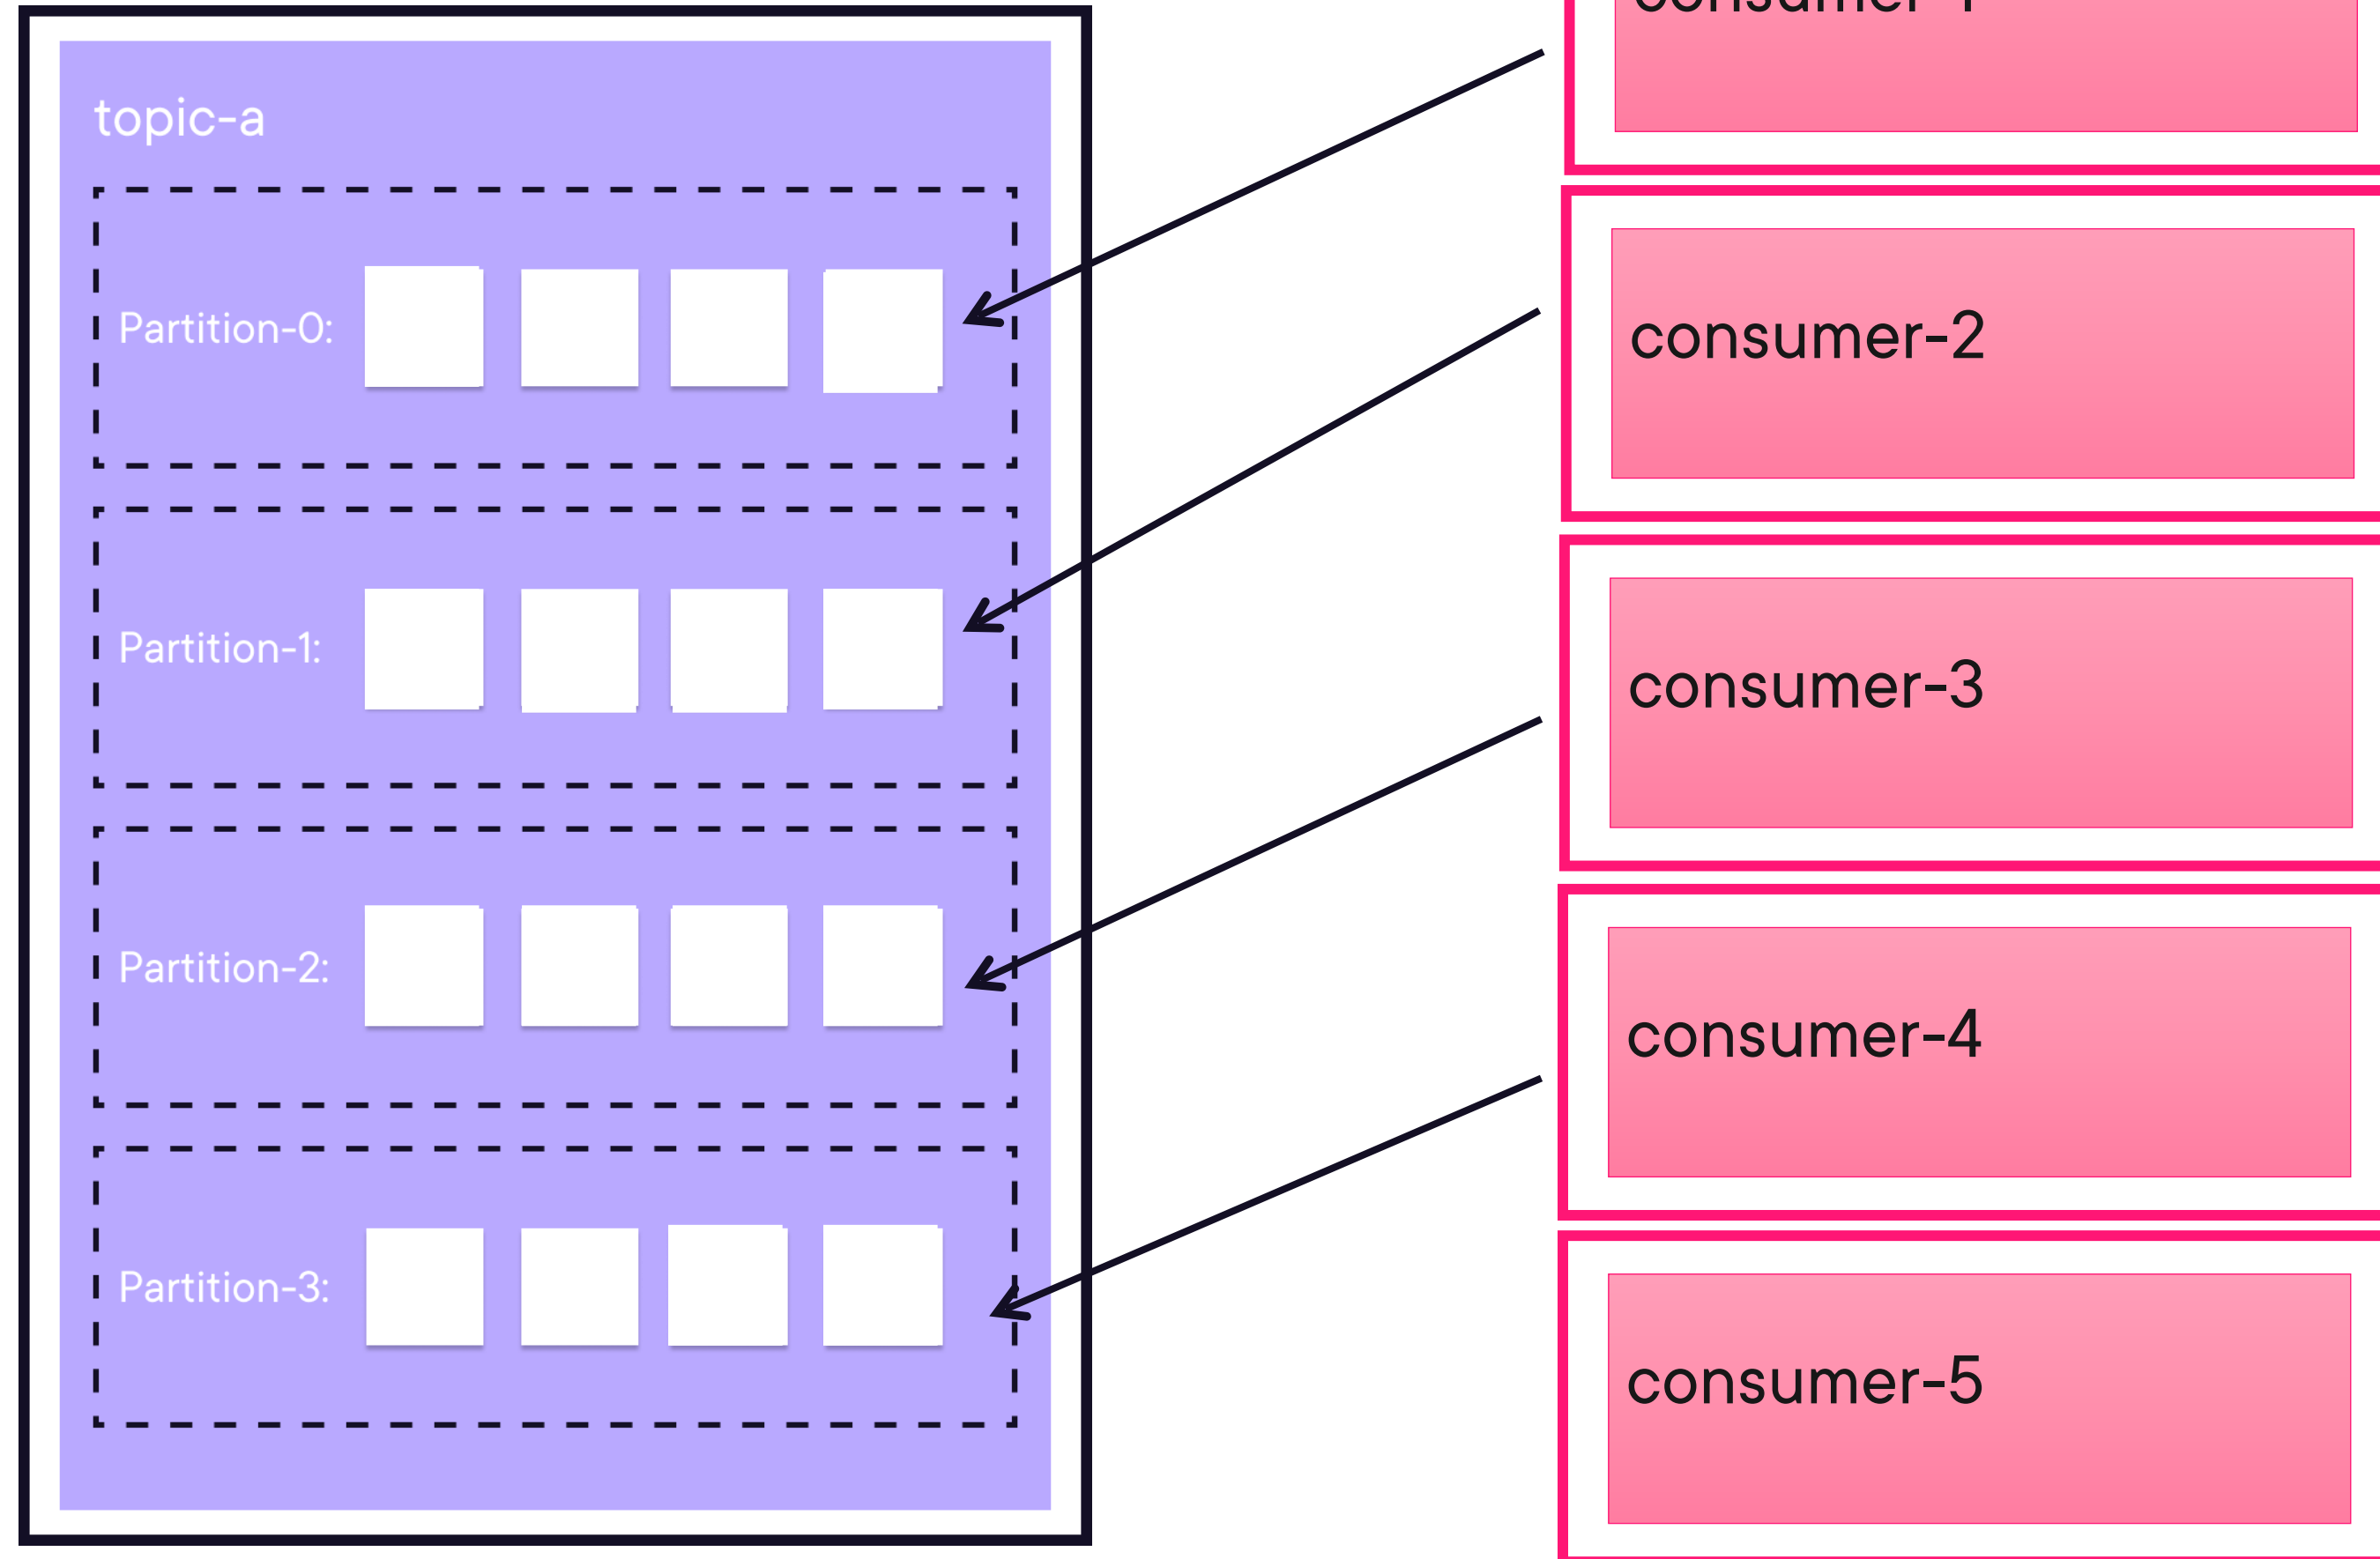


**The consumer group id is key
so Kafka knows that messages
should be distributed to both
consumers without duplicating**










Kafka Consumer Group



Consumers are typically done as a group



A single consumer will end up inefficient with large amounts of data



A consumer may never catch up



Every consumer should be on its own machine, instance, pod





Consuming from Kafka



Deserializers and Consumer Configuration



Construct a java.util.Properties object

Deserializers and Consumer Configuration

```
Properties properties = new Properties();

properties.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");

properties.put(ConsumerConfig.GROUP_ID_CONFIG, "my_group");

properties.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG, "org.apache.kafka
.common.serialization.StringDeserializer");

properties.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG, "org.apache.kaf
ka.common.serialization.IntegerDeserializer");

properties.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");

KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);
```



Provide two or more locations where the Bootstrap servers are located

Deserializers and Consumer Configuration

```
Properties properties = new Properties();

properties.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");

properties.put(ConsumerConfig.GROUP_ID_CONFIG, "my_group");

properties.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG, "org.apache.kafka
.common.serialization.StringDeserializer");

properties.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG, "org.apache.kaf
ka.common.serialization.IntegerDeserializer");

properties.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");

KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);
```



**Provide a “Team Name”,
officially called a group.id**

Deserializers and Consumer Configuration

```
Properties properties = new Properties();

properties.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");

properties.put(ConsumerConfig.GROUP_ID_CONFIG, "my_group");

properties.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG, "org.apache.kafka
.common.serialization.StringDeserializer");

properties.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG, "org.apache.kaf
ka.common.serialization.IntegerDeserializer");

properties.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");

KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);
```



Provide a DeSerializer for the key

Deserializers and Consumer Configuration

```
Properties properties = new Properties();

properties.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");

properties.put(ConsumerConfig.GROUP_ID_CONFIG, "my_group");

properties.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG, "org.apache.kafka
.common.serialization.StringDeserializer");

properties.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG, "org.apache.kaf
ka.common.serialization.IntegerDeserializer");

properties.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");

KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);
```



Provide a DeSerializer for the value

Deserializers and Consumer Configuration

```
Properties properties = new Properties();

properties.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");

properties.put(ConsumerConfig.GROUP_ID_CONFIG, "my_group");

properties.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG, "org.apache.kafka
.common.serialization.StringDeserializer");

properties.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG, "org.apache.kaf
ka.common.serialization.IntegerDeserializer");

properties.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");

KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);
```



Consumer Object

**Construct a
org.apache.kafka.clients.consumer.KafkaConsumer object**

```
KafkaConsumer consumer = new KafkaConsumer<>(properties);
```



**Use the consumer
that you have
constructed, and call
poll with pulse time.**

Processing

```
while (!done.get()) {
    ConsumerRecords<String, String> records =
        consumer.poll(Duration.of(500, ChronoUnit.MILLIS));
    for (ConsumerRecord<String, String> record : records) {
        System.out.format("offset: %d\n", record.offset());
        System.out.format("partition: %d\n", record.partition());
        System.out.format("timestamp: %d\n", record.timestamp());
        System.out.format("timeStampType: %s\n",
            record.timestampType());
        System.out.format("topic: %s\n", record.topic());
        System.out.format("key: %s\n", record.key());
        System.out.format("value: %s\n", record.value());
    }
}
```



Processing

Be a Good Citizen

```
consumer.close();
```



Speaking of Rebalances



Regardless of what a consumer is doing, at regular intervals, which are configurable, they send a heartbeat to Kafka




Specifically to the coordinator inside the broker, basically letting them know they are alive




Also when we poll for records that lets Kafka know that consumer is alive



Speaking of Rebalances




And as we said when you close a consumer that close method lets Kafka know they are not alive anymore



Which triggers a rebalance







Consuming Messages with Java





Key, Takeaways, and Tips



Takeaways



Consumers act as a group via the Consumer group ID



Each consumer in the group gets assigned a partition, and a partition is not shared by two members of a group



A rebalance is triggered when a member leaves the group or they haven't sent a heartbeat in a long time



A rebalance is a stop the world event that ensures all partitions are attended by some consumer in the group



Keys



Be sure to configure your Consumer to ensure you are not duplicating messages



Try to create a consumer loop with two consumers and verify the partition assignament playing with keys in the producer



See what happens in the logs of the broker when a consumer joins or leaves the group



Up Next:

Kafka Streams

